

# How to “Think Cloud”

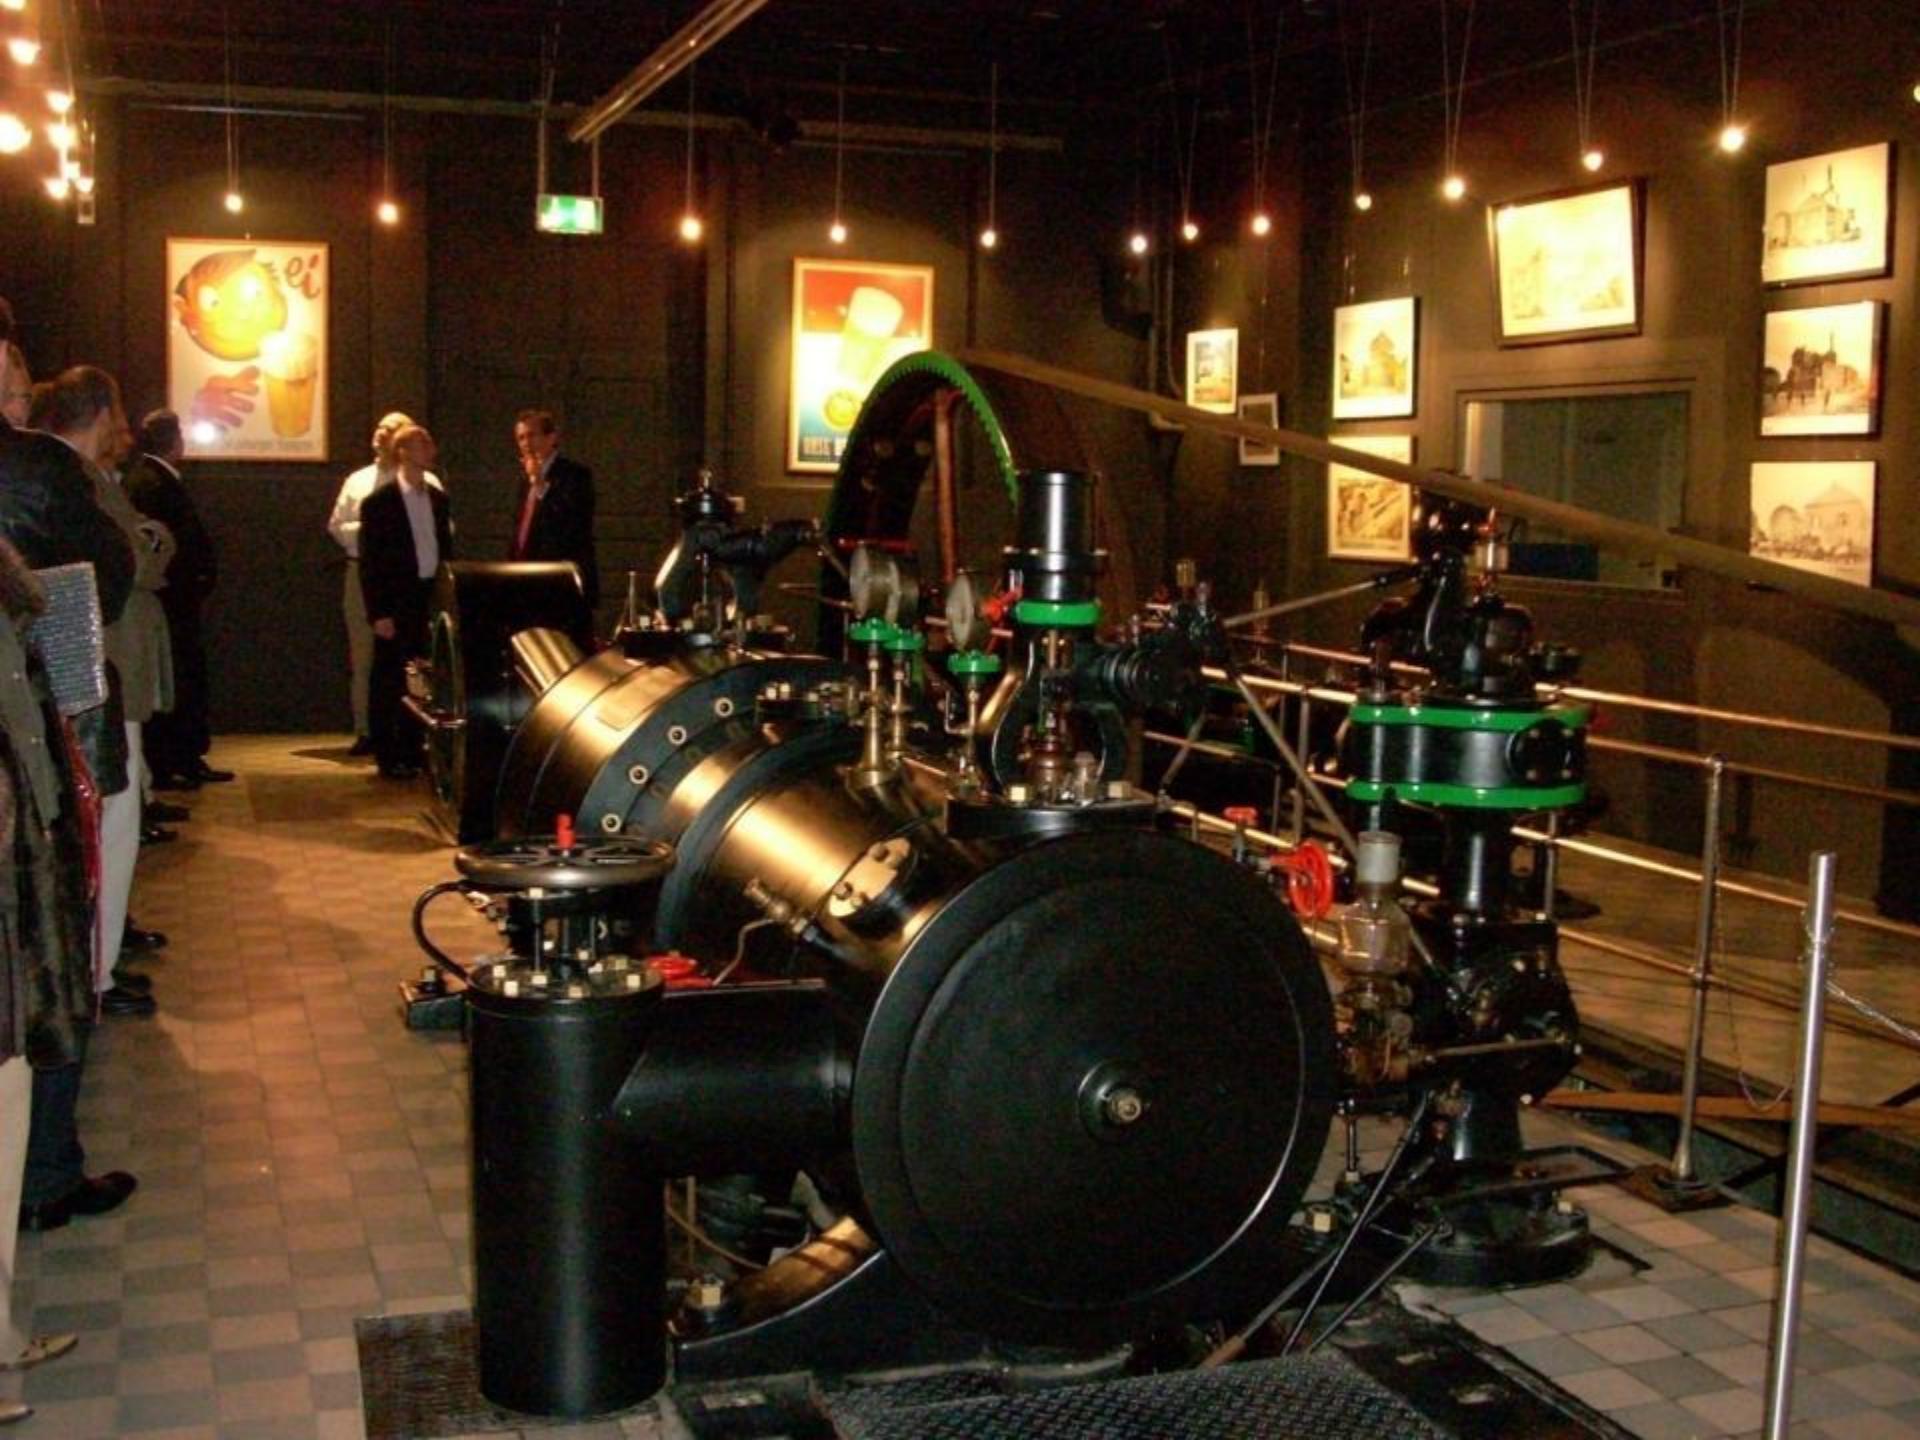
*Architectural Design  
Patterns for Cloud Computing*

# Cloud Best Practices Whitepaper

Prescriptive guidance to Cloud Architects

[http://media.amazonwebservices.com/  
AWS\\_Cloud\\_Best\\_Practices.pdf](http://media.amazonwebservices.com/AWS_Cloud_Best_Practices.pdf)



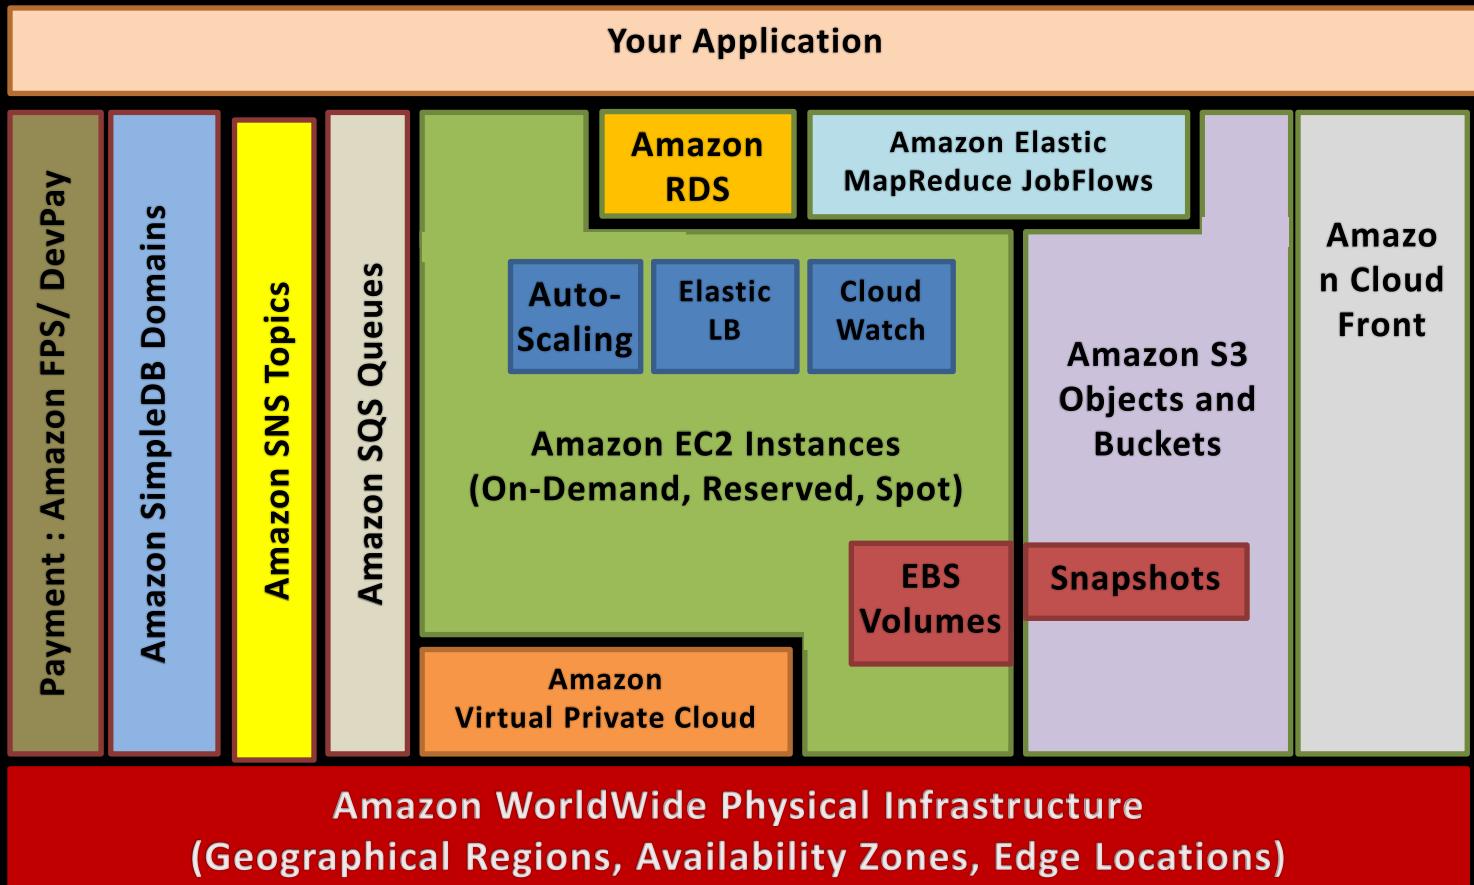


# The “Living and Evolving” Cloud

AWS services and basic terminology

## Most Applications Need:

1. Compute
2. Storage
3. Messaging
4. Payment
5. Distribution
6. Scale
7. Analytics



# Cloud Computing Attributes

What makes the Cloud so attractive

Abstract  
Resources

Focus on your needs, not on hardware specs. As your needs change, so should your resources.

On-Demand  
Provisioning

Ask for what you need, exactly when you need it. Get rid of it when you don't need

Scalability in  
minutes

Scale out or in depending on usage needs.

Pay per  
consumption

No long-term commitments.  
Pay only for what you use.

Efficiency of  
Experts

Utilize the skills, knowledge and resources of experts.

# Scalability

## Build Scalable Architecture on AWS

A scalable architecture is critical to take advantage of a scalable infrastructure

### Characteristics of Truly Scalable Service

Increasing resources results in a proportional increase in performance

A scalable service is capable of handling heterogeneity

A scalable service is operationally efficient

A scalable service is resilient

A scalable service becomes more cost effective when it grows

# Cloud Architecture Lessons

using Amazon Web Services

- 
1. Design for failure and nothing fails
  2. Loose coupling sets you free
  3. Implement “Elasticity”
  4. Build Security in every layer
  5. Don't fear constraints
  6. Think Parallel
  7. Leverage different storage options

# 1. Design for Failure

and nothing will really fail



"Everything fails, all the time"  
*Werner Vogels, CTO Amazon.com*

Avoid single points of failure

Assume everything fails, and design backwards

Goal: Applications should continue to function even if the underlying physical hardware fails or is removed or replaced.

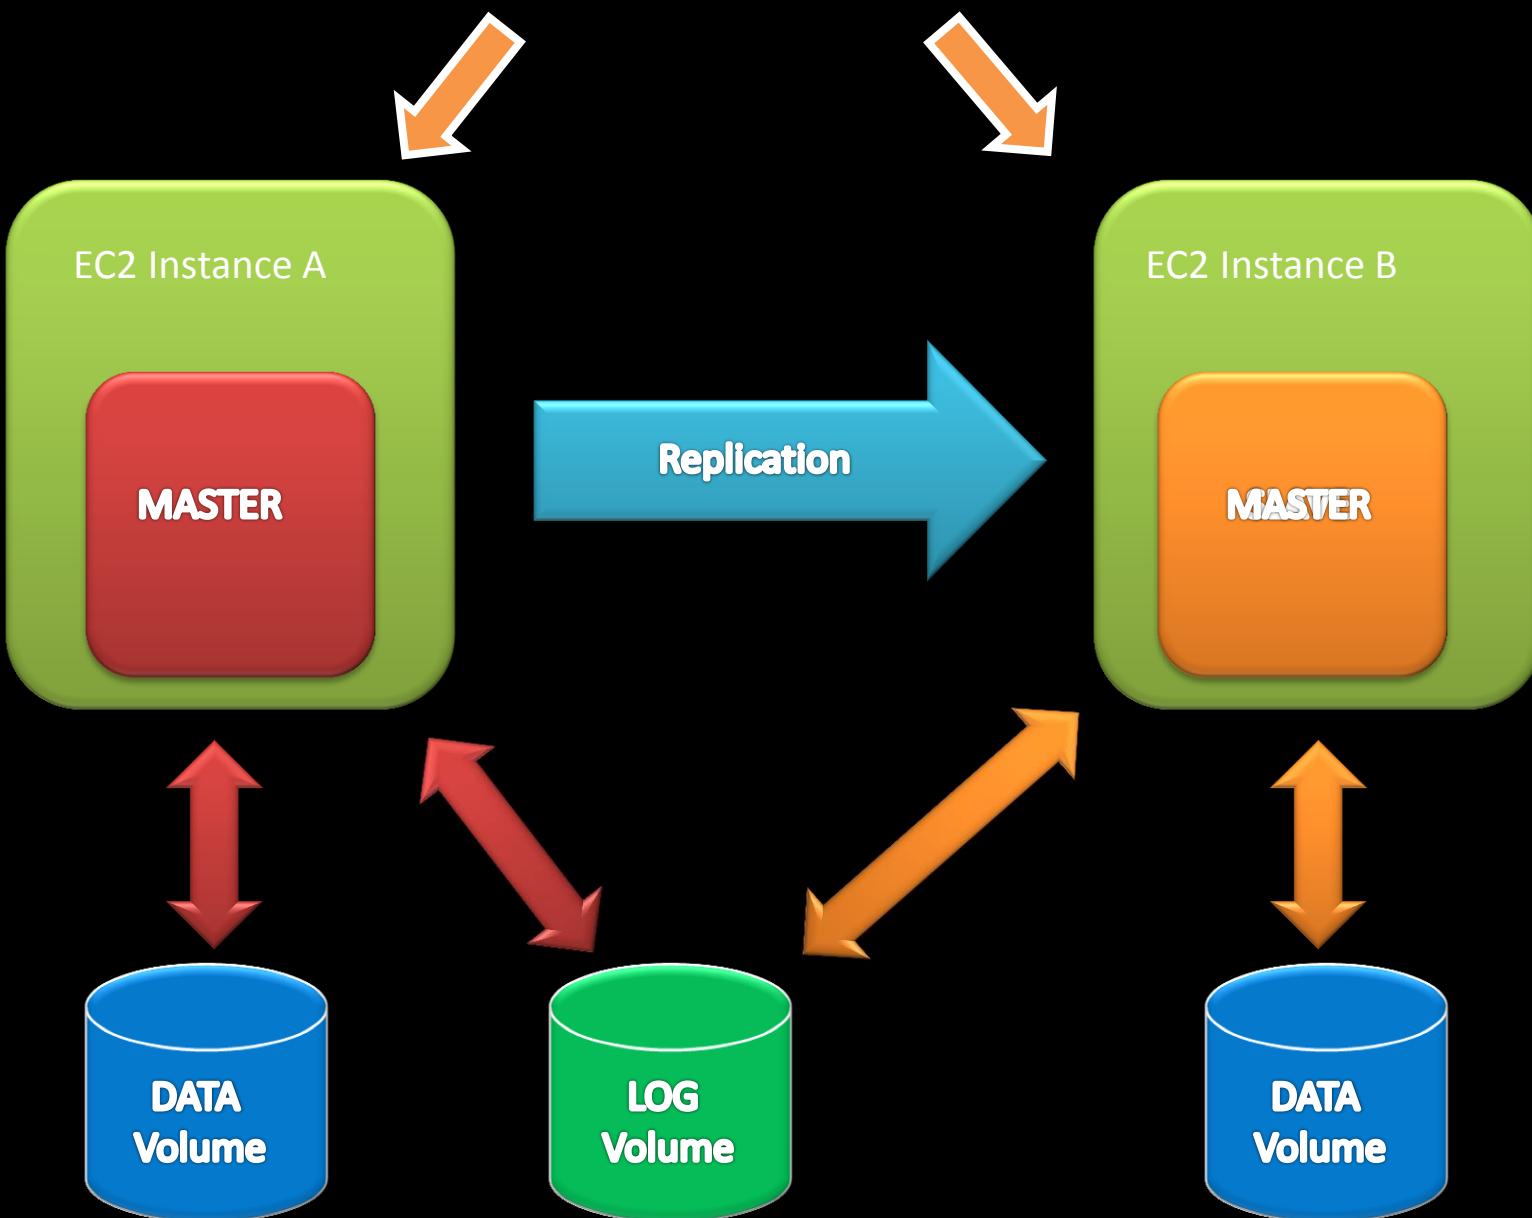
# Design for Failure with AWS

Tools to make your life easier



- Use Elastic IP addresses for consistent and re-mappable routes
- Use multiple Amazon EC2 Availability Zones (AZs)
- Create multiple database slaves across AZs
- Use real-time monitoring (Amazon CloudWatch)
- Use Amazon Elastic Block Store (EBS) for persistent file systems

YourWebsite.com



Availability Zone 1



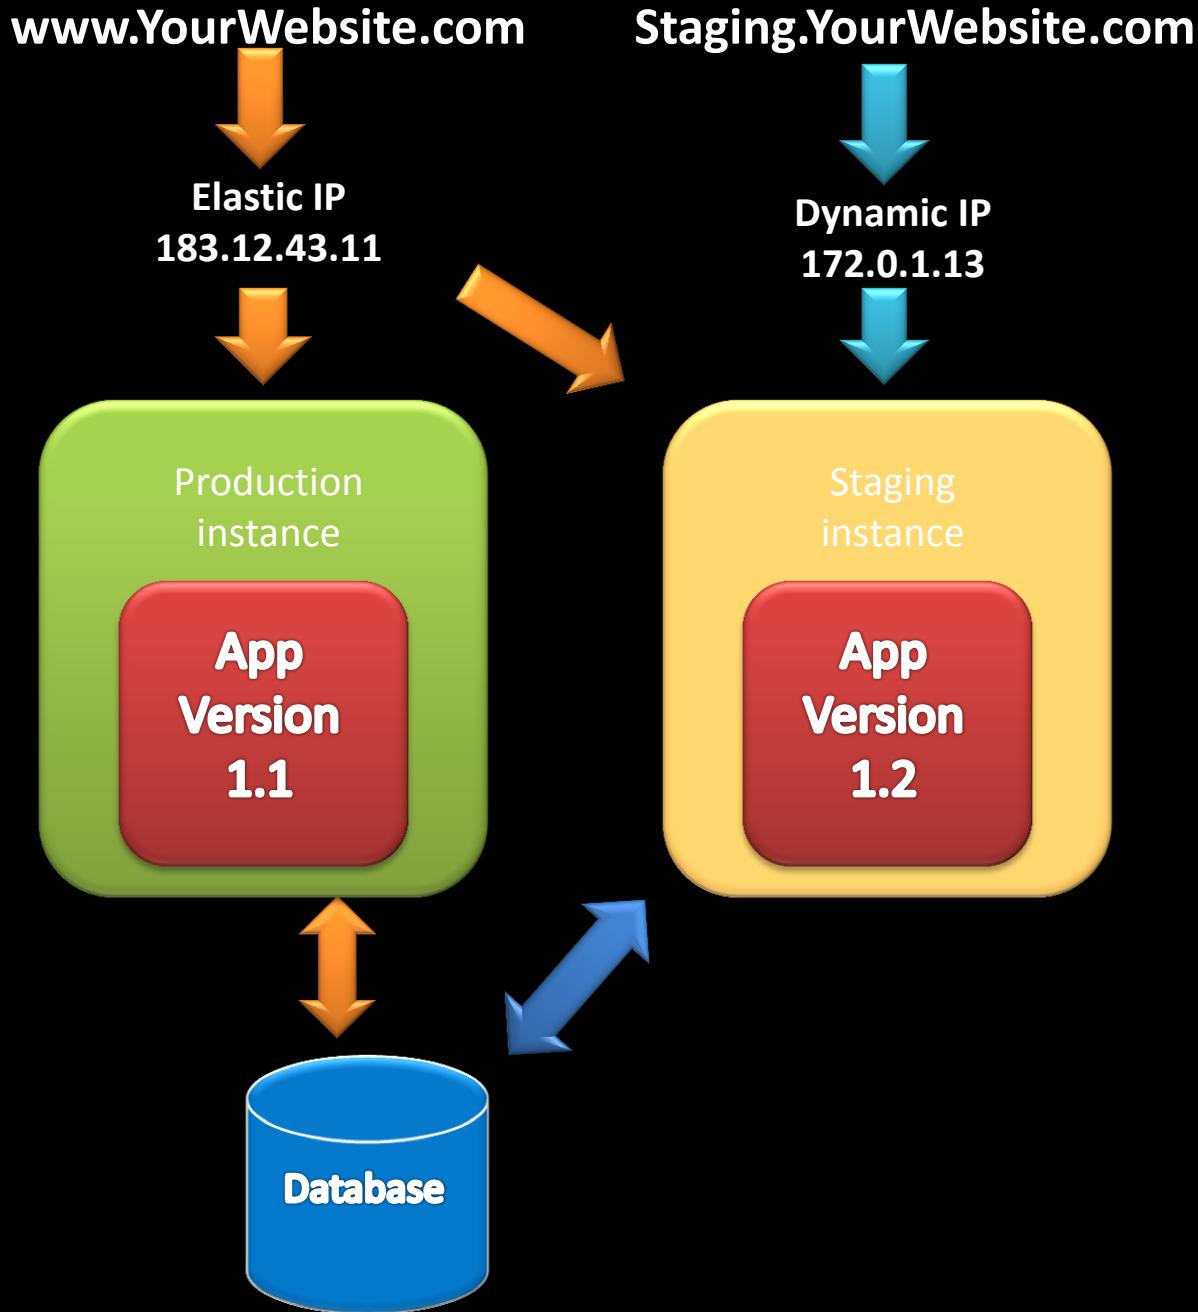
Replication



Availability Zone 2



Amazon S3



## 2. Build Loosely Coupled Systems

The looser they're coupled, the bigger they scale

Independent components

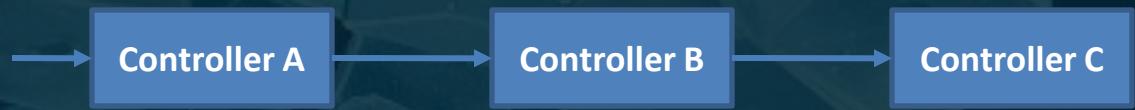
Design everything as a Black Box

De-coupling for Hybrid models

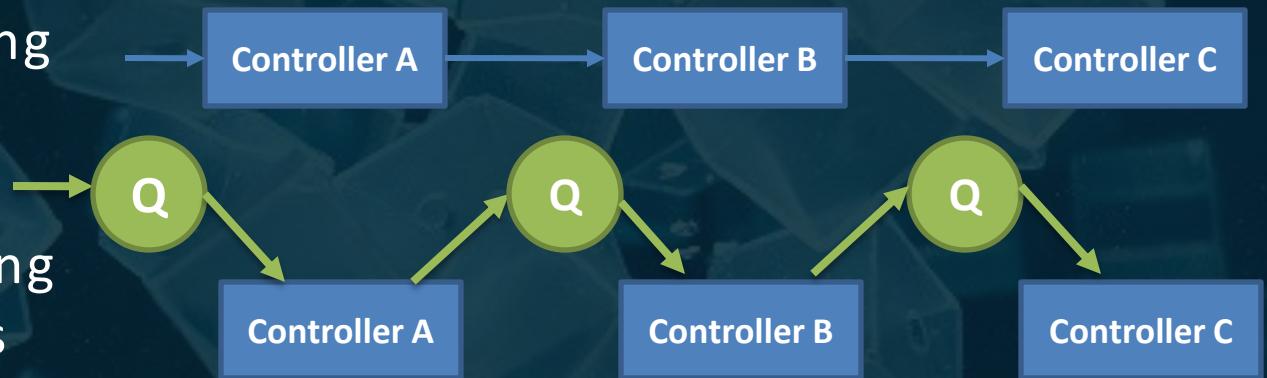
Load-balance clusters

Use Amazon SQS as Buffers

Tight Coupling



Loose Coupling  
using Queues



### 3. Implement Elasticity

Elasticity is fundamental property of the Cloud

Don't assume health or fixed location of components  
Use designs that are resilient to reboot and re-launch  
**Bootstrap** your instances: Instances on boot will ask a question "*Who am I & what is my role?*"  
Enable dynamic configuration

Use Auto-scaling (Free)

Use Elastic Load Balancing on multiple layers

Use configurations in SimpleDB to bootstrap instance

### 3. Implement Elasticity

Automate everything



Dev/Test

Apps  
Prod

Managed  
Development  
Environment

AWS Cloud

SMB IT Dept

SaaS

Paid  
AMI

Automated  
Deployment  
Environment

AWS Cloud

ISV

Web 2.0 Marketing  
Campaign

Cloud-powered  
Software Lifecycle  
management

AWS Cloud

Startup

# 3. Implement Elasticity

Standardized Application Stacks



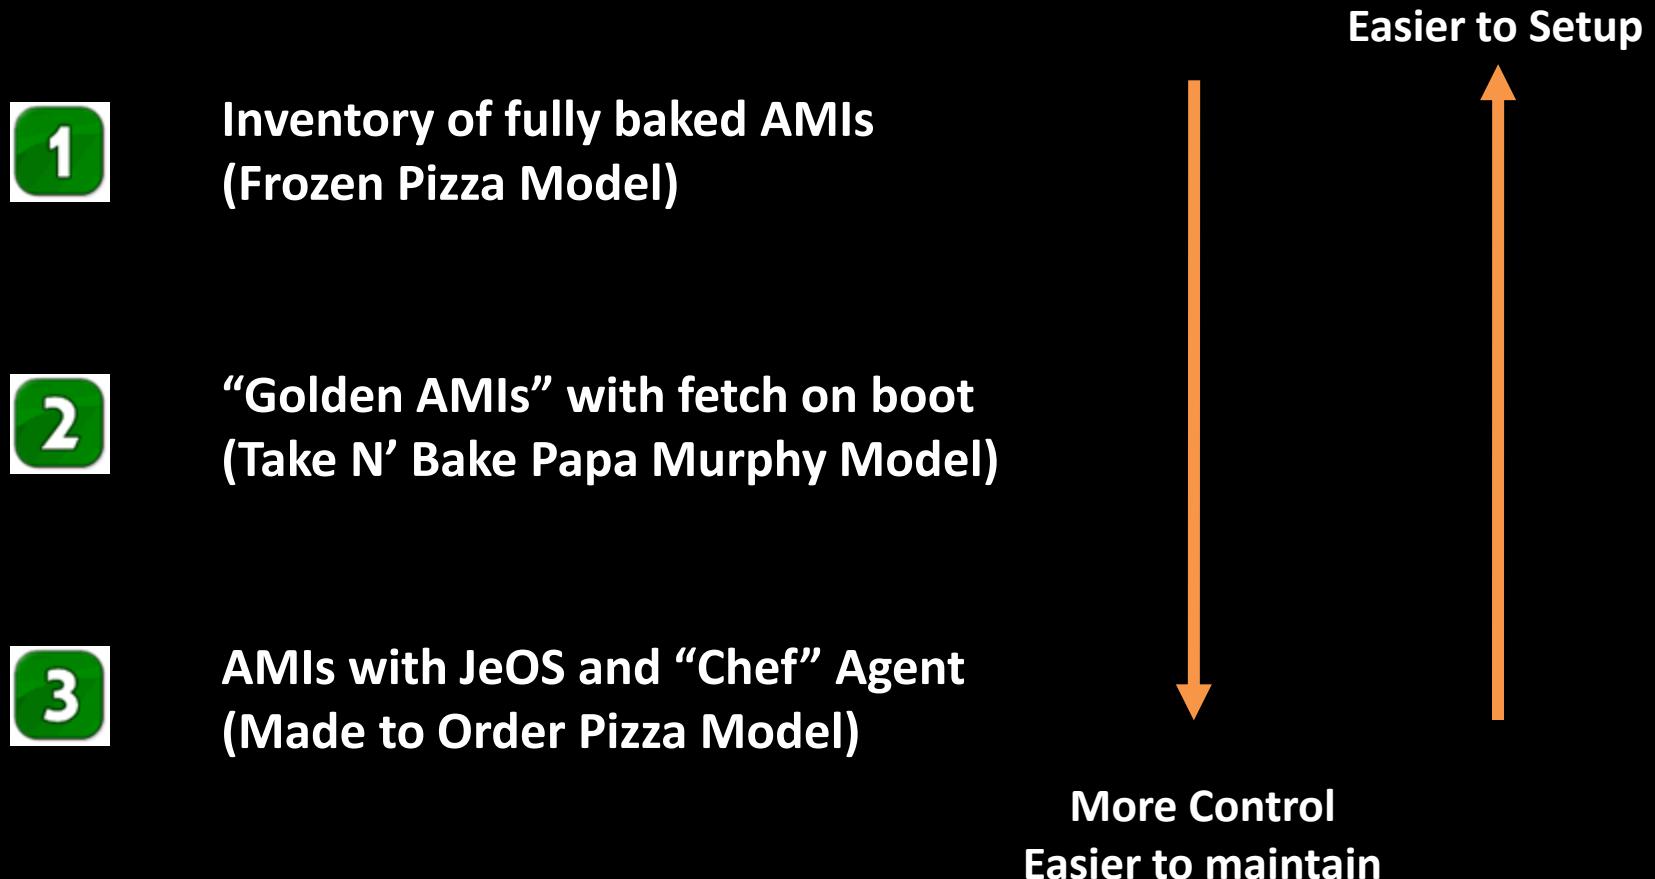
Java Stack

.NET Stack

RoR stack

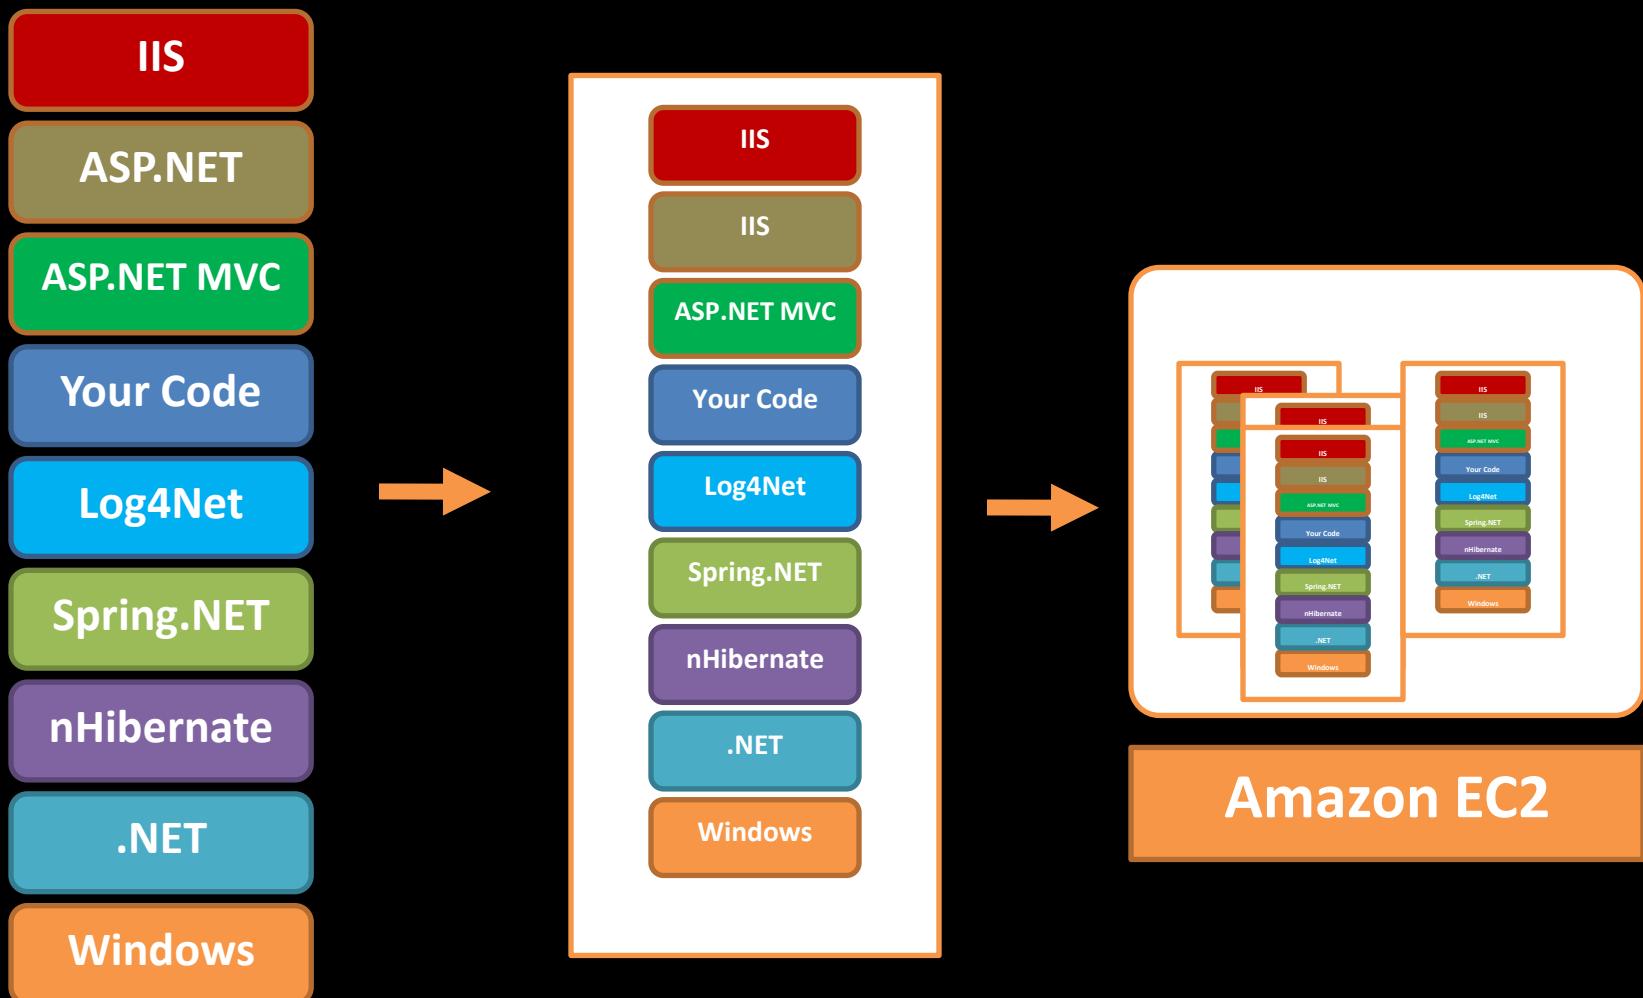
# 3. Implement Elasticity

3 approaches to designing your AMIs



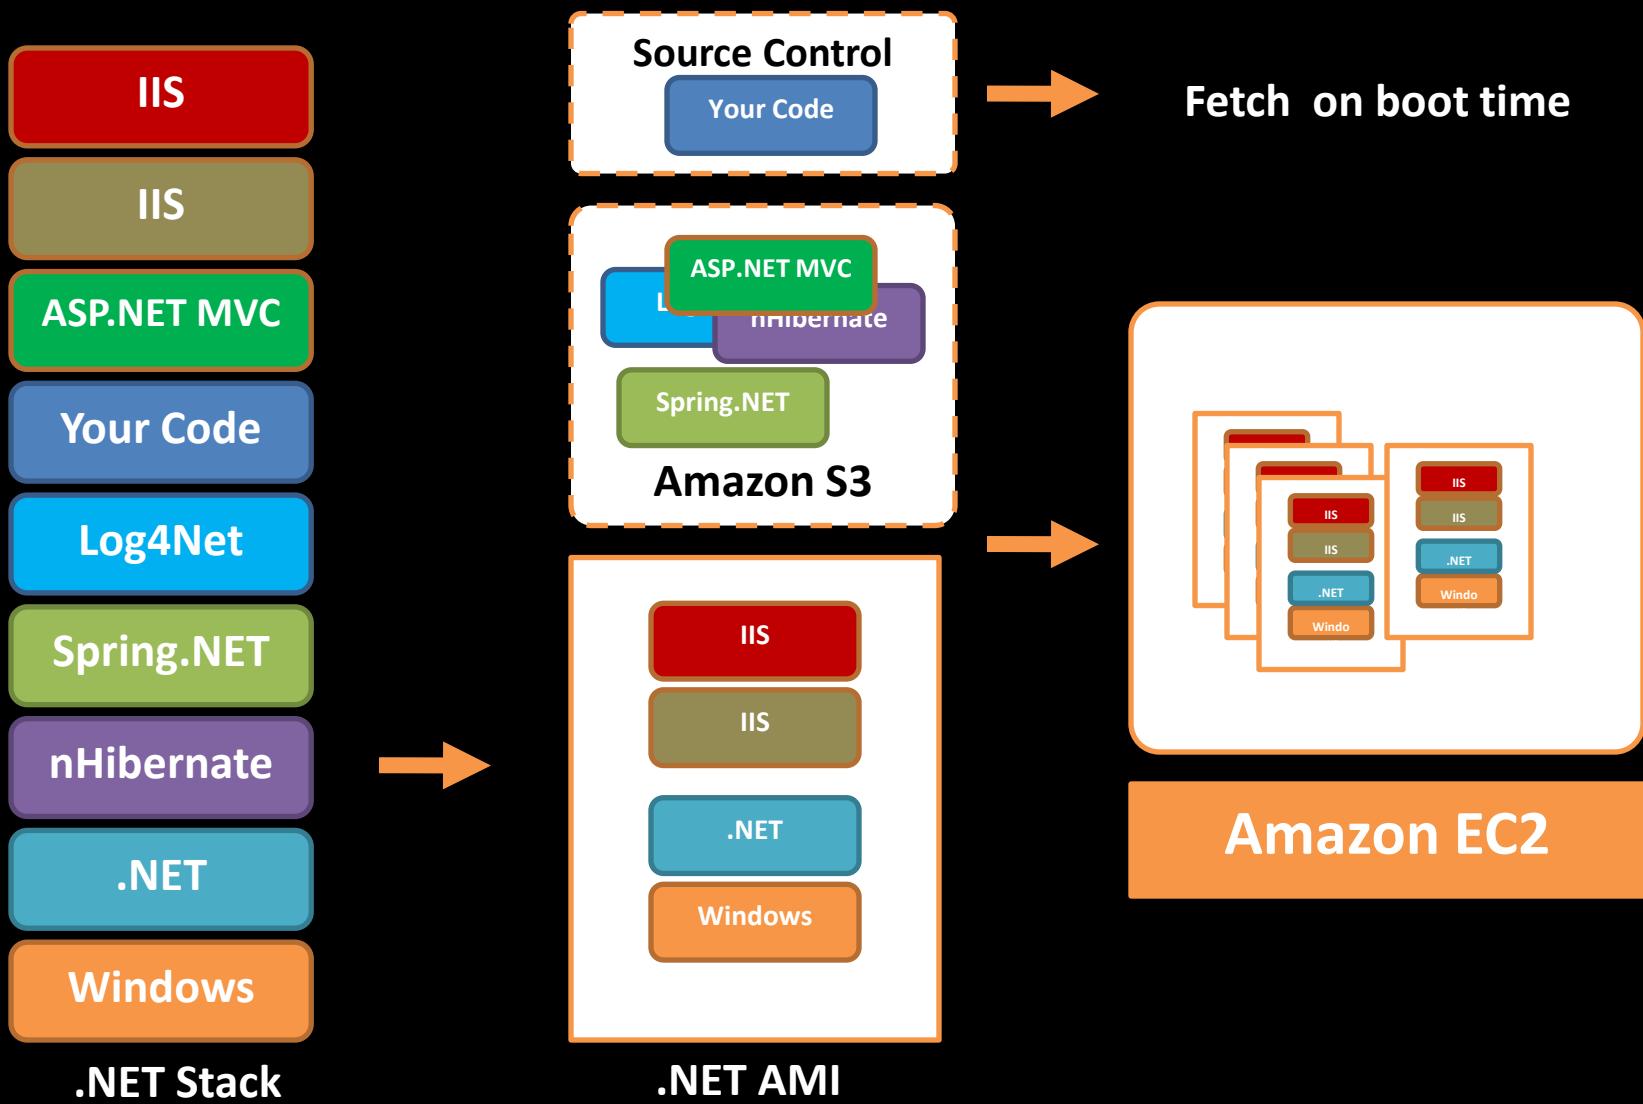
# 3. Implement Elasticity

## 1. Frozen Pizza Model



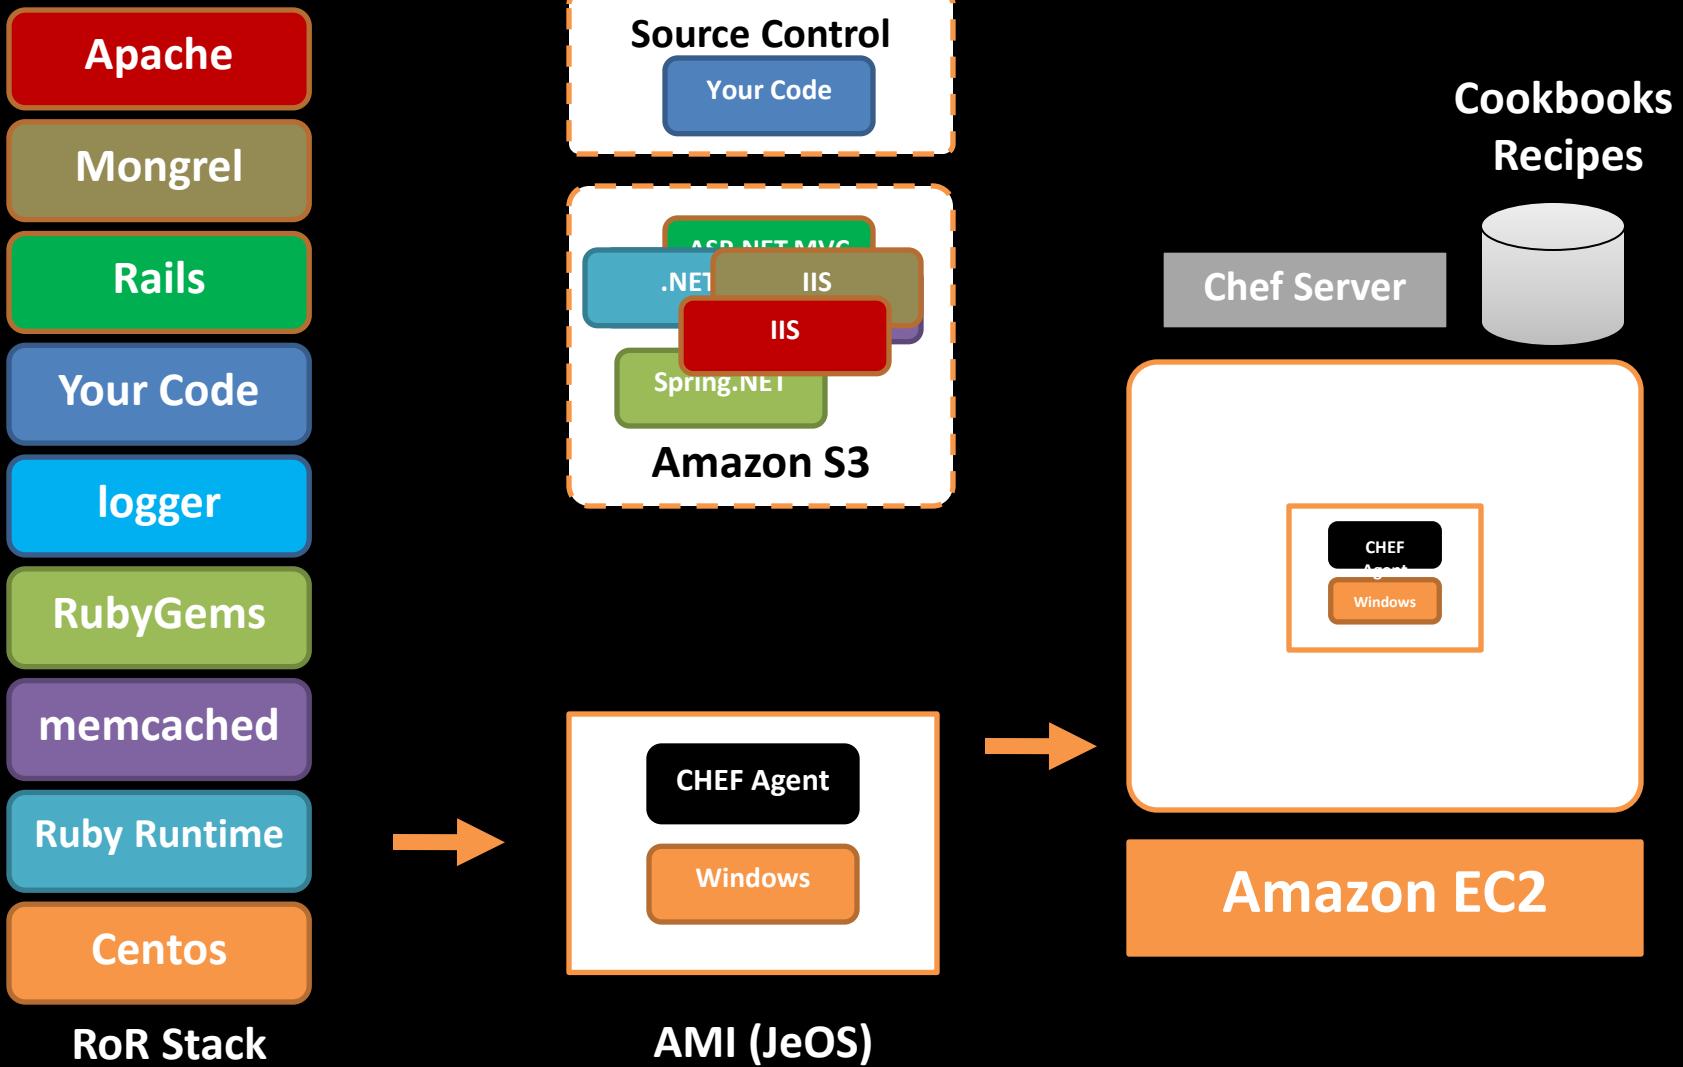
# 3. Implement Elasticity

## 2. Papa Murphy Pizza Model



# 3. Implement Elasticity

## 3. Made to Order Pizza Model

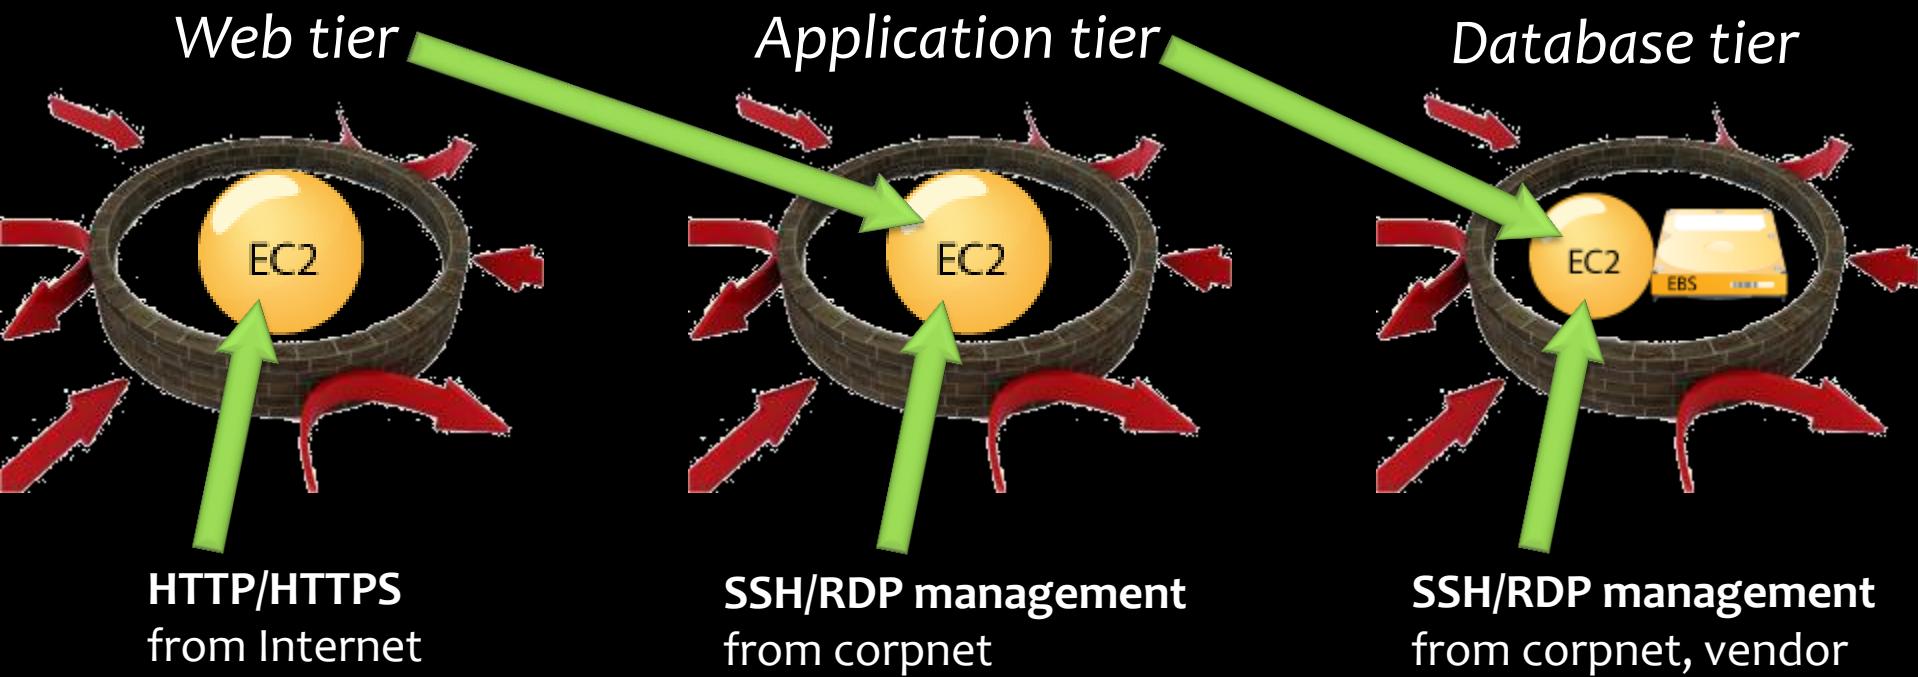


## 4. Build Security in every layer

Design with Security in mind

With cloud, you lose a little bit of physical control but not your ownership

- Create distinct Security Groups for each Amazon EC2 cluster
- Use group-based rules for controlling access between layers
- Restrict external access to specific IP ranges
- Encrypt data “at-rest” in Amazon S3
- Encrypt data “in-transit” (SSL)
- Consider encrypted file systems in EC2 for sensitive data
- Rotate your AWS Credentials, Pass in as arguments encrypted
- Use MultiFactor Authentication

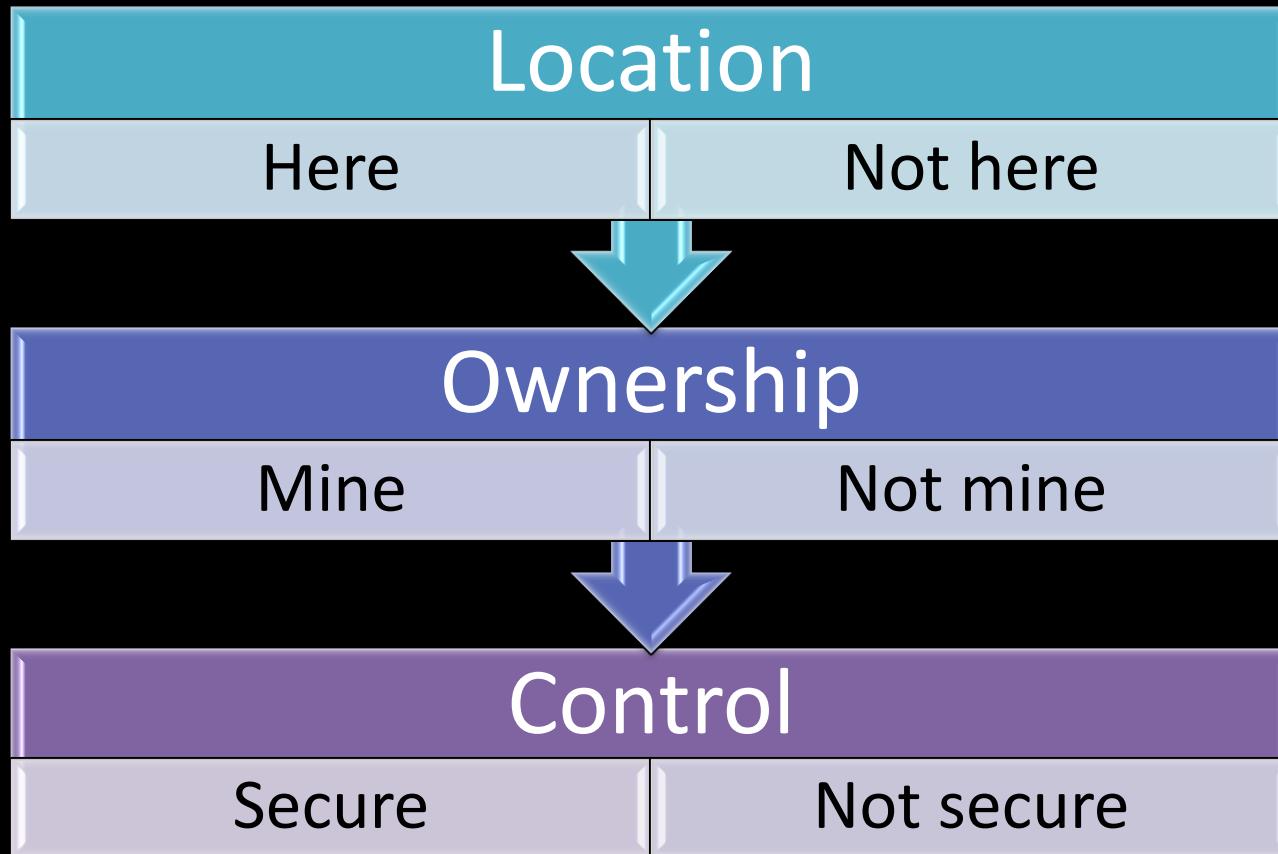


```
ec2-authorize WebSG -P tcp -p 80 -s 0.0.0.0/0
ec2-authorize WebSG -P tcp -p 443 -s 0.0.0.0/0

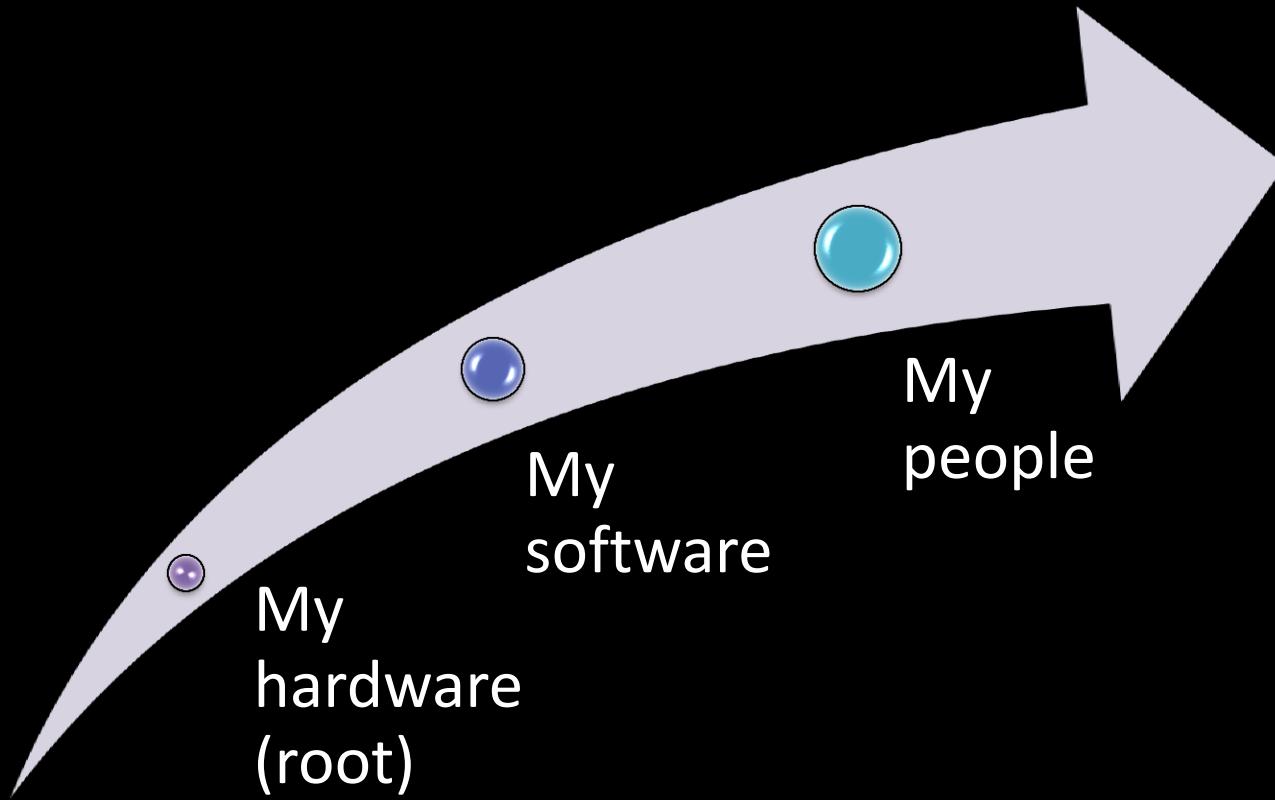
ec2-authorize AppSG -P tcp -p AppPort -o WebSG
ec2-authorize AppSG -P tcp -p 22|3389 -s CorpNet

ec2-authorize DBSG -P tcp -p DBPort -o AppSG
ec2-authorize DBSG -P tcp -p 22|3389 -s CorpNet
ec2-authorize DBSG -P tcp -p 22|3389 -s Vendor
```

# Traditional security model

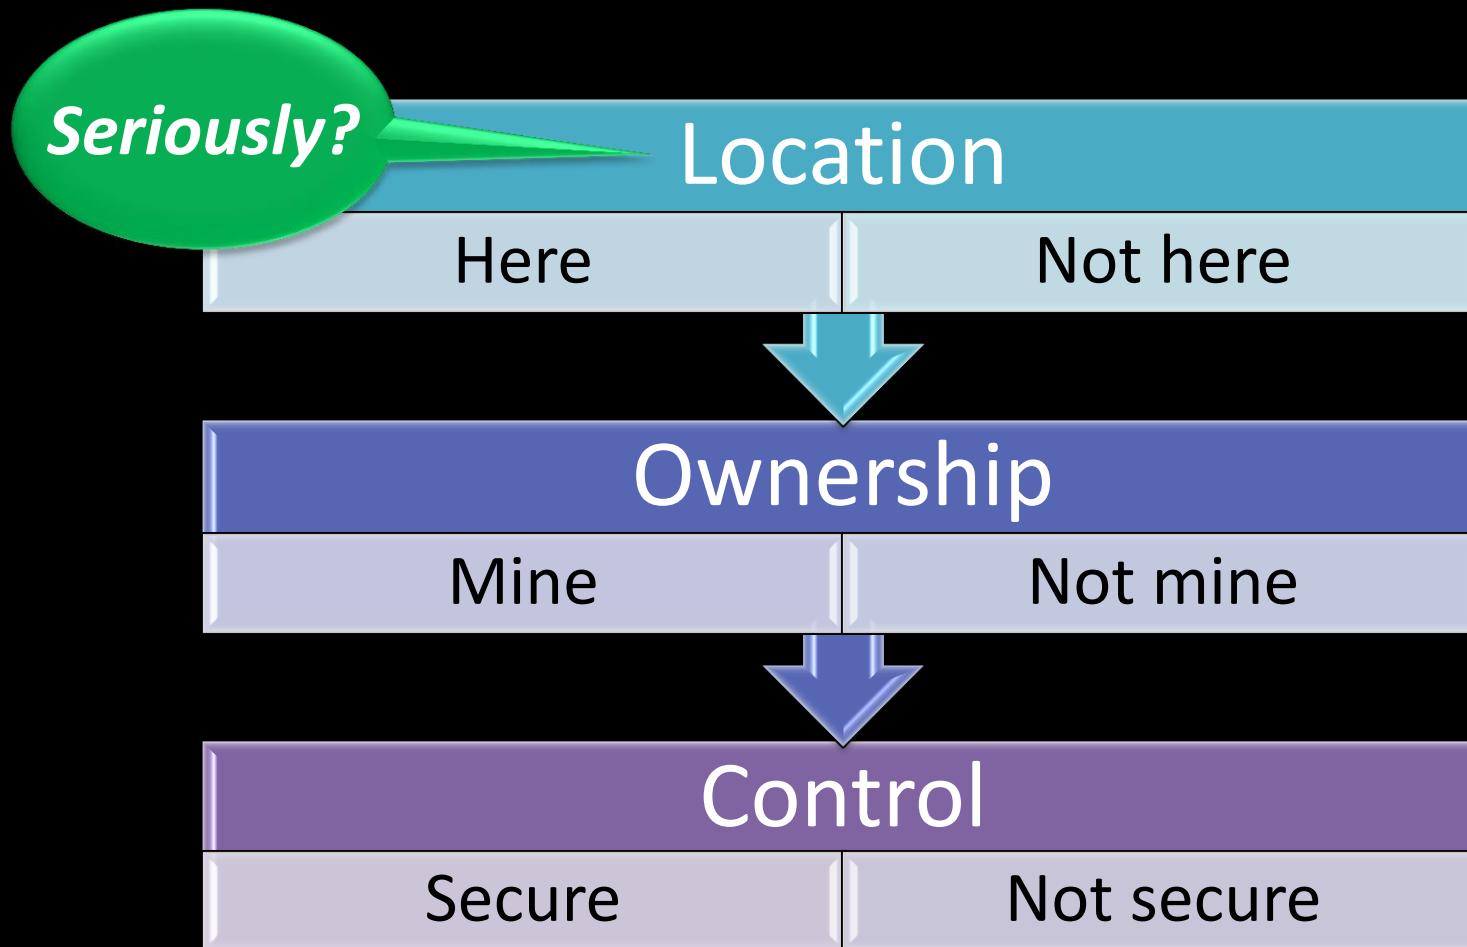


# Layers of trust

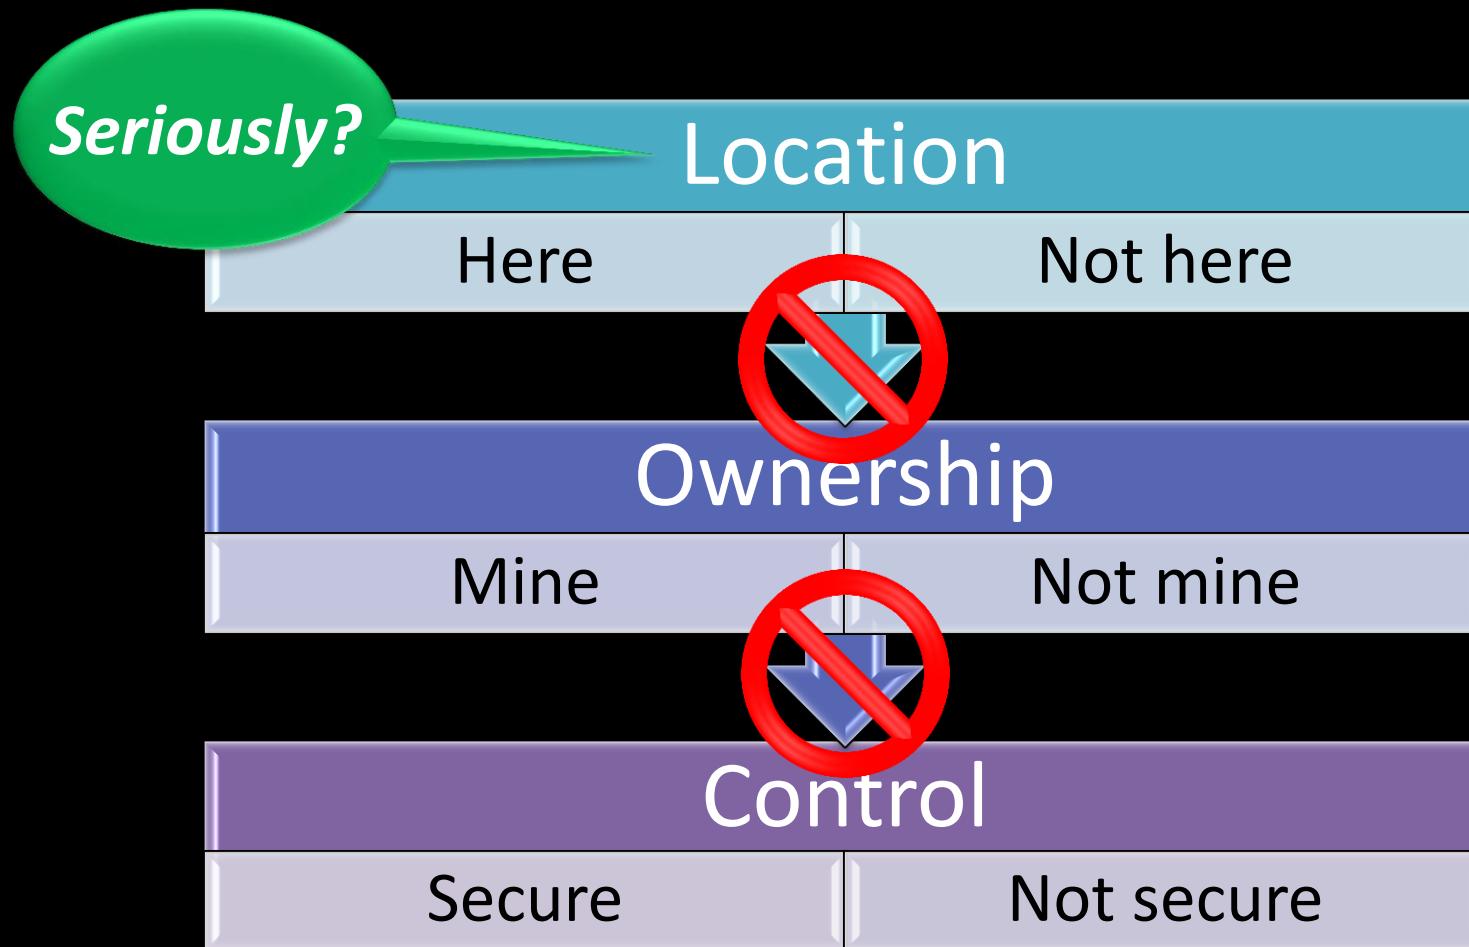


Perimeters separate trusted (owned, local)  
from untrusted (other, remote)

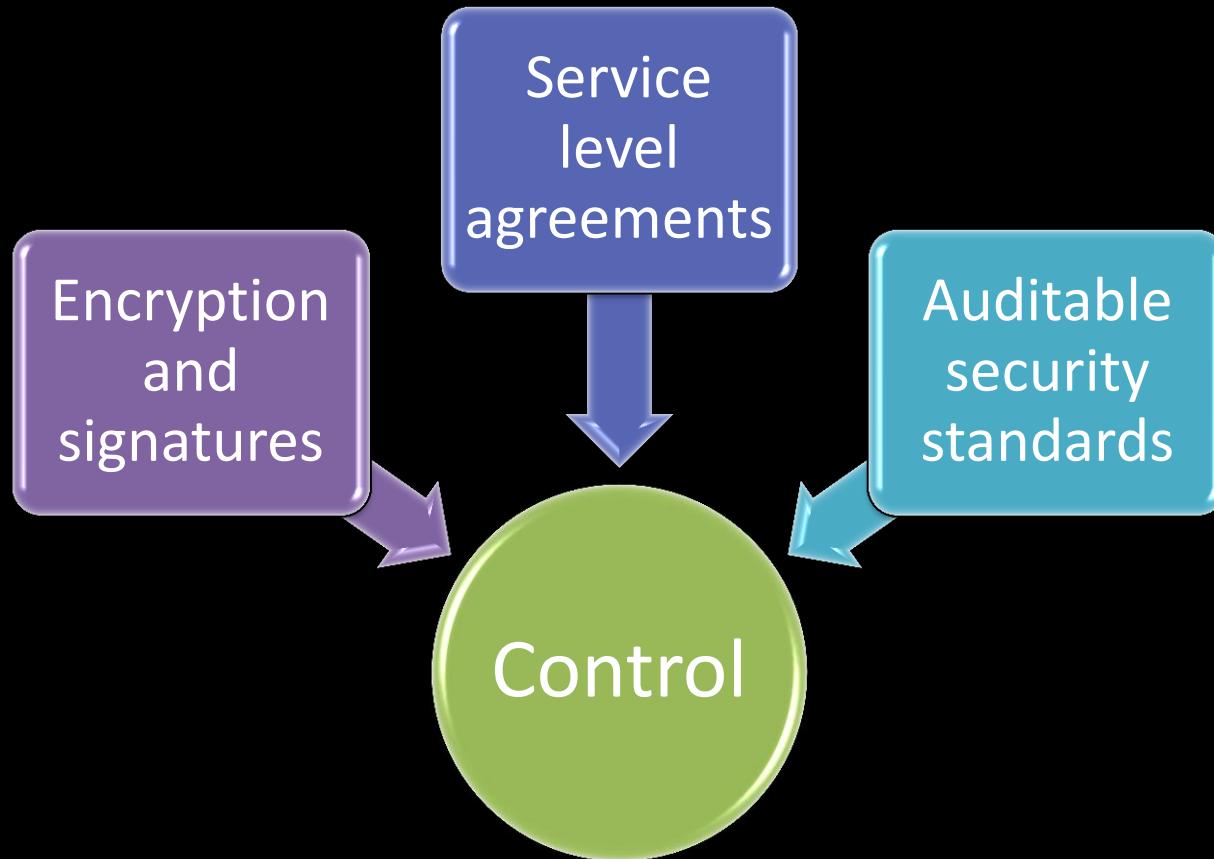
# The model is breaking



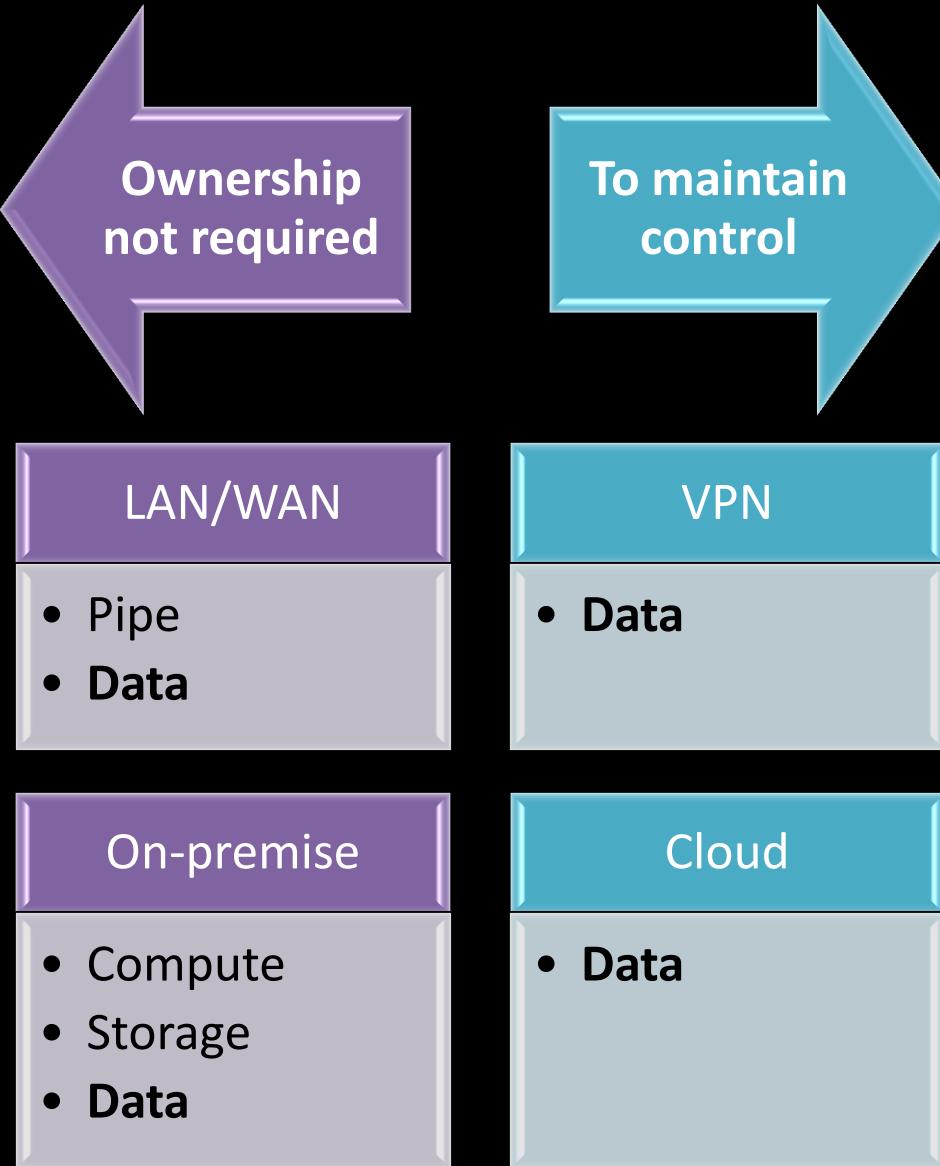
# The model is breaking



# New security model



# Ownership vs. control



## 5. Don't fear constraints

Re-think architectural constraints

More RAM? Distribute load across machines  
Shared distributed cache

Better IOPS on my database?

Multiple read-only / sharding / DB clustering

Your hardware failed or messed up config?  
simply throw it away and switch to new hardware with no additional cost

Hardware Config does not match?  
Implement Elasticity

Performance

Caching at different levels (Page, Render, DB)

# 6. Think Parallel

Serial and Sequential is now history

Experiment different architectures in parallel

Multi-threading and Concurrent requests to cloud services

Run parallel MapReduce Jobs

Use Elastic Load Balancing to distribute load across multiple servers

Decompose a Job into its simplest form – and with “shared nothing”

The beauty of the cloud shines when you combine elasticity and parallelization

## 6. Leverage many storage options

One size DOES NOT fit all



Amazon S3: large static objects

Amazon Cloudfront: content distribution

Amazon SimpleDB: simple data indexing/querying

Amazon EC2 local disc drive : transient data

Amazon EBS: persistent storage for any RDBMS + Snapshots on S3

Amazon RDS: RDBMS service - Automated and Managed MySQL

# 6. Leverage many storage options

Which storage option to use when?

	Amazon S3 + CF	Amazon EC2 Ephemeral Store	Amazon EBS	Amazon SimpleDB	Amazon RDS
Ideal for	Storing Large write-once, read-many types of objects, Static Content Distribution	Storing non-persistent transient updates	Off-instance persistent storage for any kind of data,	Querying light-weight attribute data	Storing and querying structured Relational and referential Data
Ideal examples	Media files, audio, video, images, Backups, archives, versioning	Config Data, scratch files, TempDB	Clusters, boot data, Log or data of commercial RDBMS like Oracle, DB2	Querying, Mapping, tagging, click-stream logs, metadata, shared-state management, indexing	Complex transactional systems, inventory management and order fulfillment systems
Not recommended for	Querying, Searching	Storing Database logs or backups, customer data		Relational (joins) query	
Not recommended examples	Database, File Systems	Sensitive data	Content Distribution	OLTP, DW cube rollups	Simple lookups

# Cloud Architecture Lessons

## Best Practices

1. Design for failure and nothing fails
2. Loose coupling sets you free
3. Implement Elasticity
4. Build Security in every layer
5. Don't fear constraints
6. Think Parallel
7. Leverage many storage options



# Migrating your Web Application

Step by Step towards AWS

A typical Web App needs:

- Compute Power
- Storage capacity
- Content Distribution
- Database storage
- Messaging
- Load balancing
- Monitoring

With AWS:

- Amazon EC2
- Amazon S3
- Amazon CloudFront
- Amazon EBS
- Amazon SQS
- Amazon EC2
- Amazon CloudWatch

# Amazon Web Services tools

Things you need

A black and white photograph of various wrenches and tools laid out on a textured, metallic surface. The tools include several open-end wrenches of different sizes, a crescent wrench, and other small hand tools. A dark, semi-transparent rounded rectangle is overlaid on the lower half of the image, containing the following text:

**Web** : AWS Management Console  
**IDE** : AWS Toolkit for Eclipse  
**AWS SDK**: .NET SDK, Java SDK  
**Tools** : 3<sup>rd</sup> Party tools eg. CA  
**Firefox Plugins** :  
ElasticFox, S3Fox, SDB Tool  
**Several libraries**: boto, cloudfusion

# Identify the right candidate

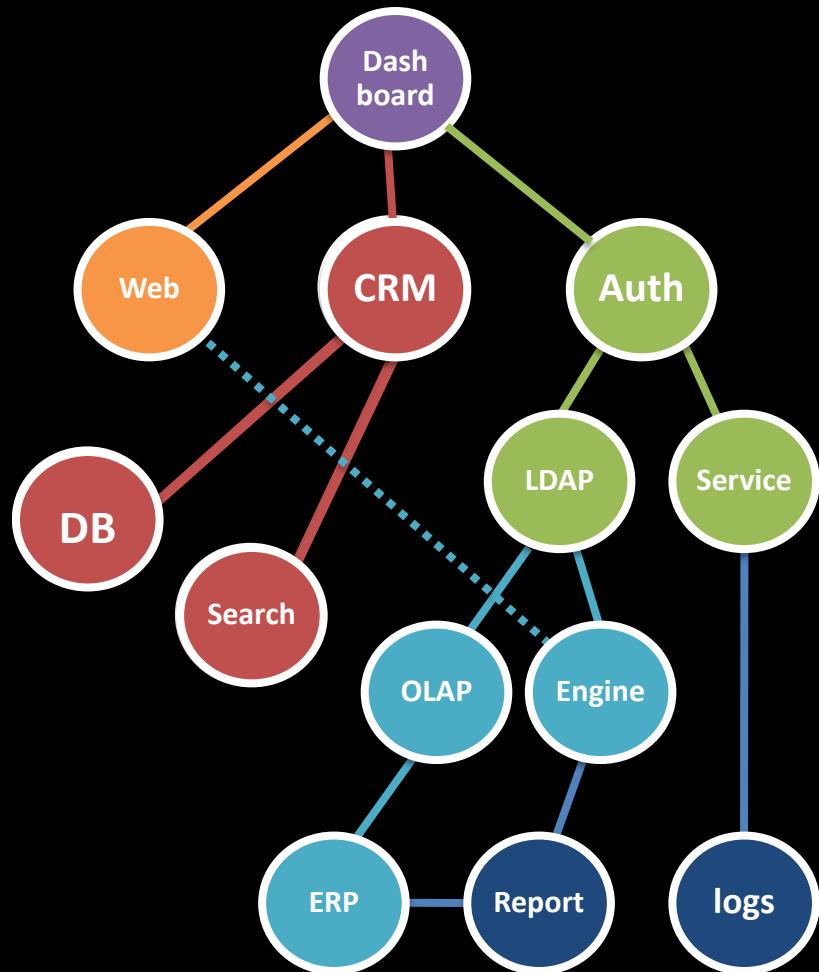
## Assessment



List all your IT assets  
Whiteboard your IT Assets  
Identify upward and downward dependencies

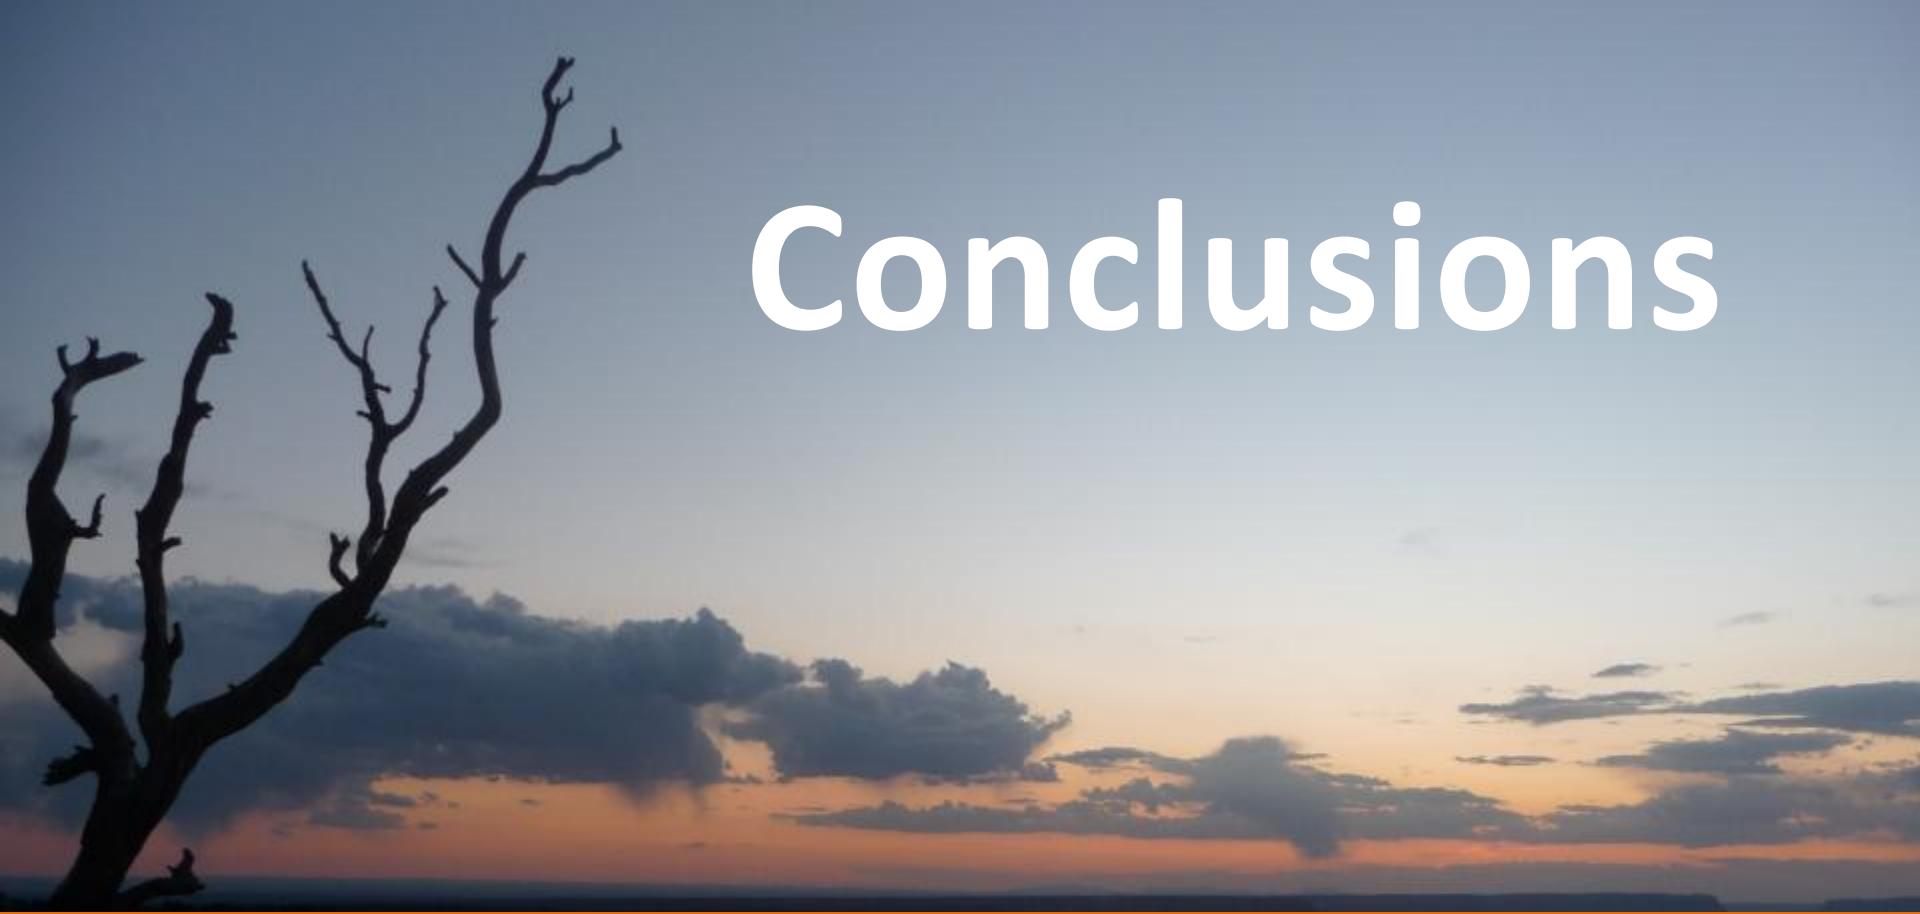
# Identify the right candidate

Pick one application with lower dependencies to start with



Search for under-utilized IT assets  
Applications that has immediate business need to scale  
Applications that are running out of capacity

**Low-hanging fruits (Examples):**  
Web Applications  
Batch Processing systems  
Build/QA/Test systems  
Content Management Systems  
Digital Asset Management Systems



# Conclusions

## **Most Important Lesson From Our Customers:**

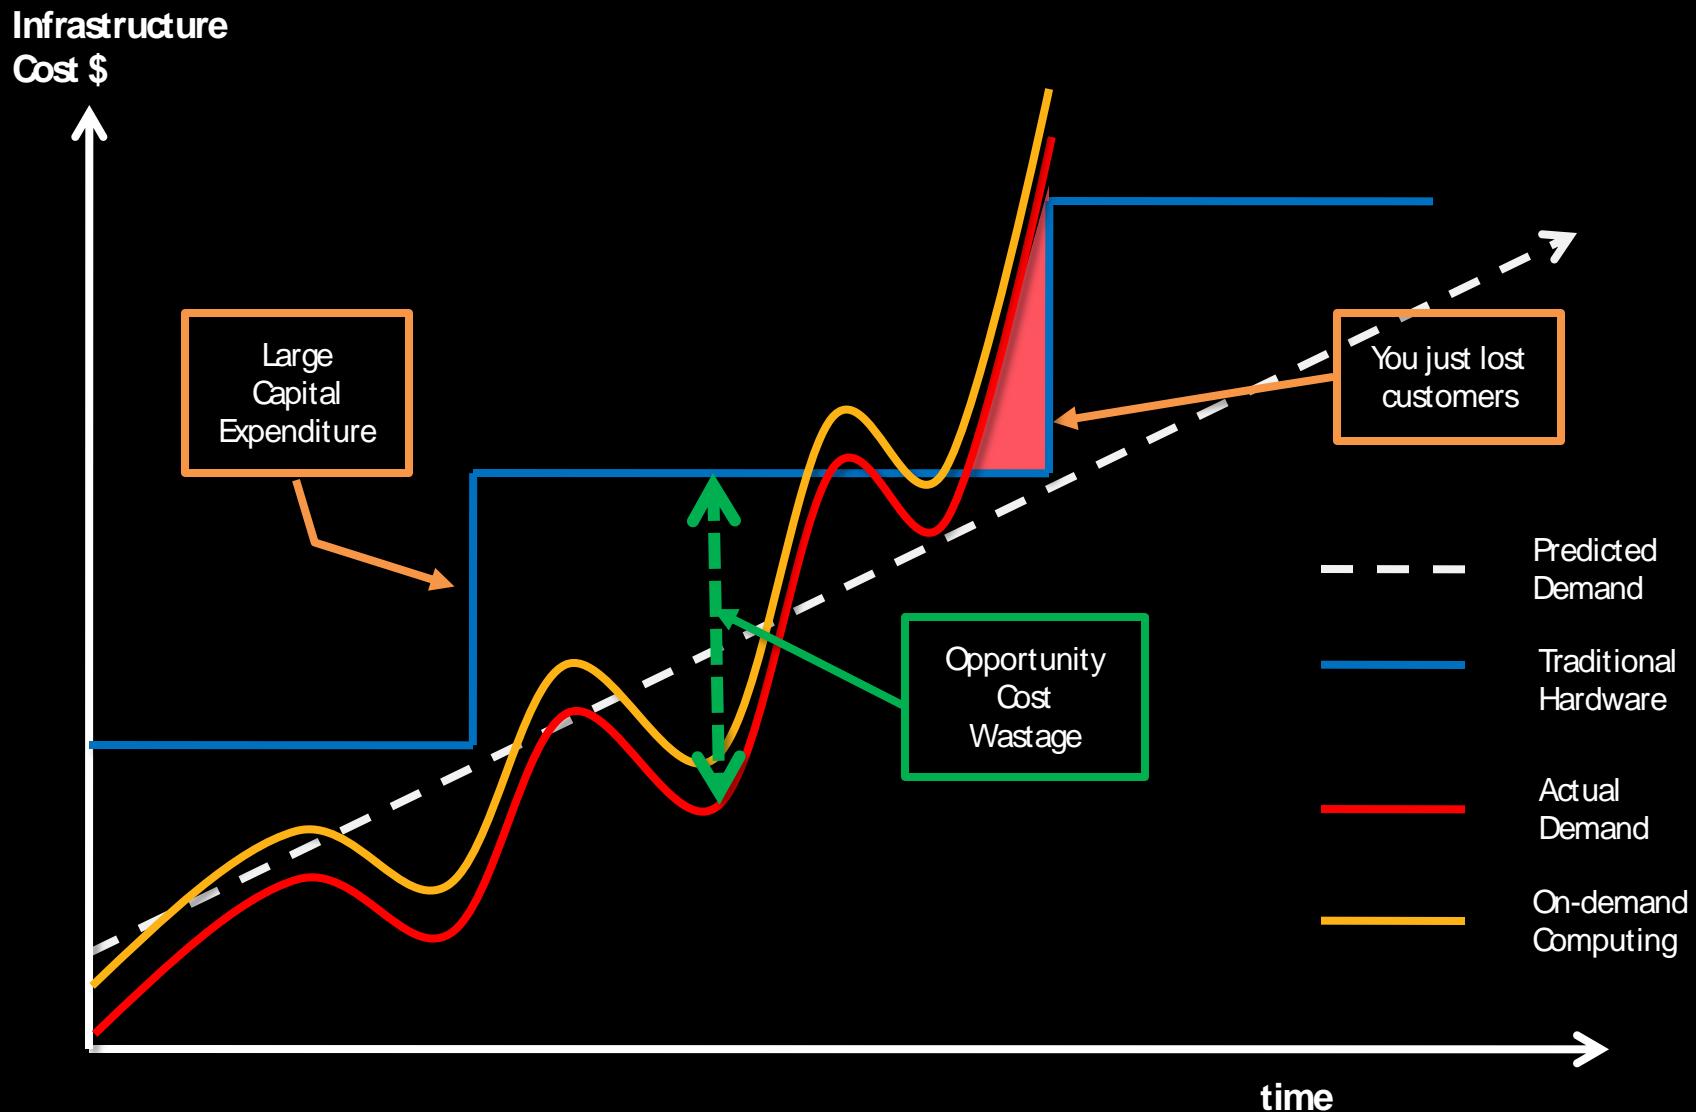
Start small with a well-defined proof of concept

Experiment with different architectures; Keep one, throw away others

Once one application is launched others will follow...

## **Traditional IT roles are changing**

# Predicting Infrastructure Needs



# The day is not too far....

Scalability, Security, High availability, Fault-tolerance, Testability and Elasticity will be configurable properties of the application architecture and will be an automated and intrinsic part of the platform on which they are built.



# Thank you!

steriley@amazon.com  
@steveriley @awscloud  
<http://stvrlly.wordpress.com>

*Presentation ideas and template from @simon and @jinman*